

# Heaps

*Dr. Aiman Hanna*

**Department of Computer Science & Software Engineering  
Concordia University, Montreal, Canada**

**These slides have been extracted, modified and updated from original slides of :**

**Data Structures and Algorithms in Java, 5th edition. John Wiley & Sons, 2010. ISBN 978-0-470-38326-1.**

**Data Structures and the Java Collections Framework by William J. Collins, 3rd edition, ISBN 978-0-470-48267-4.**

**Both books are published by Wiley.**

**Copyright © 2010-2011 Wiley**

**Copyright © 2010 Michael T. Goodrich, Roberto Tamassia**


**Copyright © 2011 William J. Collins**

**Copyright © 2011-2021 Aiman Hanna**

**All rights reserved**

# Coverage & Warning

## □ Heaps



**Warning:** As an upfront warning, the “heap” data structure discussed here has nothing to do with the memory heap used in the run-time environment.

# Recall Priority Queue ADT

- A priority queue stores a collection of entries
- Each **entry** is a pair (key, value)
- Main methods of the Priority Queue ADT
  - **insert(k, x)**  
inserts an entry with key k and value x
  - **removeMin()**  
removes and returns the entry with smallest key
- Additional methods
  - **min()**  
returns, but does not remove, an entry with smallest key
  - **size(), isEmpty()**
- Applications:
  - Standby flyers
  - Auctions
  - Stock market

# Recall P.Q. Sorting



- We use a priority queue
  - Insert the elements with a series of **insert** operations
  - Remove the elements in sorted order with a series of **removeMin** operations
- The running time depends on the priority queue implementation:
  - Unsorted sequence gives selection-sort:  **$O(n^2)$**  time
  - Sorted sequence gives insertion-sort:  **$O(n^2)$**  time
- **Can we do better?**

## Algorithm ***PQ-Sort***( $S, C$ )

**Input** sequence  $S$ , comparator  $C$  for the elements of  $S$

**Output** sequence  $S$  sorted in increasing order according to  $C$

$P \leftarrow$  priority queue with comparator  $C$

**while**  $\neg S.isEmpty()$

$e \leftarrow S.remove(S.first())$

$P.insertItem(e, \emptyset)$

**while**  $\neg P.isEmpty()$

$e \leftarrow P.removeMin().getKey()$


$S.addLast(e)$

# Heaps

- Both insertion-sort and selection-sort of P.Q. achieved running time of  $O(n^2)$ .
- An efficient realization of a priority queue uses a data structure, called *heap*.
- This data structure allows us to perform both insertions and removals in logarithmic time, which is significant improvement over the list-based implementation.
- Fundamentally, the heap achieves such improvement by abandoning the idea of storing entries in a list; instead it stores the entries in a binary tree.

# Heaps

- In other words, a heap is a binary tree storing entries at its nodes.
- Additionally, a heap satisfies the following properties:
- **Heap-Order:** for every internal node  $v$  other than the root,  $key(v) \geq key(parent(v))$
- **Complete Binary Tree:** let  $h$  be the height of the heap
  - for  $i = 0, \dots, h - 1$ , there are  $2^i$  nodes of depth  $i$
  - at depth  $h - 1$ , the internal nodes are to the left of the external nodes
  - there is at most one node with a single child (that is, you cannot find two or more nodes with single child) and this child must be a left child
- The **last node** of a heap is the rightmost node of maximum depth



# Heap-order Property

- **Heap-Order Property:** This is a relational property. For every internal node  $v$  other than the root,  $key(v) \geq key(parent(v))$ .
- Consequently, the keys encountered on a path from the root to an external node are in non-decreasing order.
- Additionally, the minimum key (which is the most important one) is hence always stored at the root (or the “top of the heap”, hence the name “heap” of this data structure).

# Complete Binary Tree Property

- **Complete Binary Tree Property:** This is a structural property. This property is needed to insure that the height of the heap is as small as possible.




# Complete Binary Tree Property

- A heap  $T$  with height  $h$  is a complete binary tree if each depth  $i = 0, \dots, h - 1$  has the maximum possible number of entries, and at least one entry at the last depth (that is depth  $h$ ).
- Formally, each level  $i = 0, \dots, h - 1$  must have  $2^i$  nodes.
- In a complete binary tree, a node  $v$  is to the left of node  $w$  if  $v$  and  $w$  are at the same level and that  $v$  is encountered before  $w$  (from left to right, which is also an inorder traversal).

# Complete Binary Tree Property

- For instance, node with entry  $(15, K)$  is to the left of node with entry  $(7, Q)$ . Similarly, node with entry  $(14, E)$  is to the left of node with entry  $(8, W)$ , etc.
- Another important node is the *last node*. This is the **right-most, deepest external** node in the tree.

Note: **right-most** means the one to the right of all other nodes in the level (this may not necessary be a right child).




# Height of a Heap



- **Theorem:** A heap storing  $n$  keys has height  $O(\log n)$


Proof: (we apply the complete binary tree property)

- Let  $h$  be the height of a heap storing  $n$  keys
- Since there are  $2^i$  keys at depth  $i = 0, \dots, h-1$  and at least one key at depth  $h$ , we have  $n \geq 1 + 2 + 4 + \dots + 2^{h-1} + 1 = 2^h - 1 + 1$
- Thus,  $n \geq 2^h$ , i.e.,  $h \leq \log n$
- This theorem is very important since it implies that if we can perform updates on a heap in a time proportional to its height, then those operations will run in logarithmic time.




# Heaps and Priority Queues

- We can use a heap to implement a priority queue.
- We store a (key, element) item at each node.
- We keep track of the position of the last node.




# Insertion into a Heap

- Method  $insert(k, x)$  of the priority queue ADT corresponds to the insertion of a key  $k$  to the heap.




- The insertion algorithm consists of three steps:
  - Find the insertion node  $z$  (the new last node)
  - Store  $k$  at  $z$
  - Restore the heap-order property (discussed next)




# Upheap

- After the insertion of a new key  $k$ , the heap-order property may be violated.
- Algorithm *upheap* restores the heap-order property by swapping  $k$  along an upward path from the insertion node.





# Upheap

- Upheap terminates when the key  $k$  reaches the root or a node whose parent has a key smaller than or equal to  $k$ .
- Since a heap has height  $O(\log n)$ , upheap runs in  $O(\log n)$  time.



# Upheap

□ Example:





Heaps



# Upheap


- Example (Continues ...):



Heaps


# Upheap

- Example (Continues ...):




# Removal from a Heap (§ 7.3.3)

- Method *removeMin()* of the priority queue ADT corresponds to the removal of the root key from the heap.





- The removal algorithm consists of four steps:
  - Return the root entry
  - Replace the root key (entry in fact) with the key of the last node  $w$
  - Remove  $w$
  - Restore the heap-order property (discussed next)





# Downheap

- ❑ After replacing the root key with the key  $k$  of the last node, the heap-order property may be violated.
- ❑ Algorithm *downheap* restores the heap-order property by swapping key  $k$  along a downward path from the root.
  - If  $T$  has no right child, then the swapping (if needed) starts at the left child (which is the only one actually)
  - If both left and right children are there, then the swapping occurs with the one that has the smaller key; otherwise only one side of the tree may be corrected, which may force further swap operations.




# Downheap

- ❑ Downheap terminates when key  $k$  reaches a leaf or a node whose children have keys greater than or equal to  $k$ .
- ❑ The downward swapping process is called *down-heap bubbling*.
- ❑ Since a heap has height  $O(\log n)$ , downheap runs in  $O(\log n)$  time.



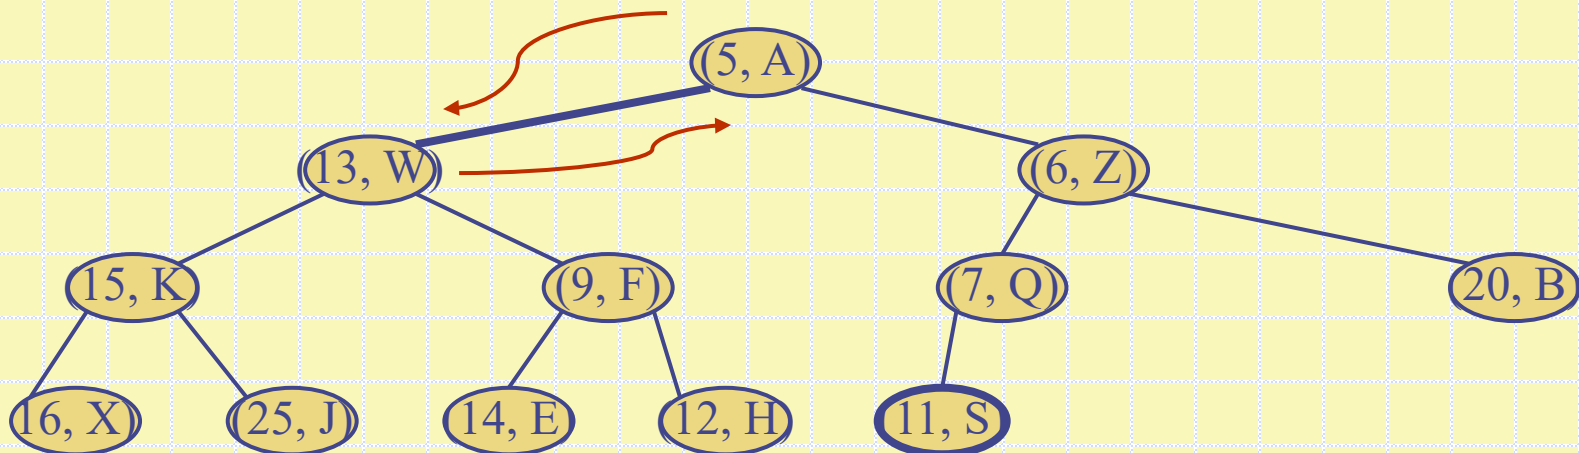

# Downheap

- Removal Example:





# Downheap

- Removal Example  
(Continues ...):



# Downheap


- Removal Example  
(Continues ...):





# Updating the Last Node

- ❑ The insertion node can be found by traversing a path of  $O(\log n)$  nodes
- ❑ For instance, if the current last node is a left child then:
  - Go up to the father and down right to the new last node
- ❑ If the current last node is a right child then:
  - Go up until a left child or the root is reached
  - If a left child is reached, go to the right child
  - Go down left until a leaf is reached
- ❑ Similar algorithm for updating the last node after a removal




# Heap Performance

- The performance of a P.Q. realized by means of a heap is as follows:

Operation	Complexity
size(), isEmpty()	$O(1)$
min()	$O(1)$
insert()	$O(\log n)$
removeMin()	$O(\log n)$

# Heap-Sort



- Let us recall the PQ-Sort algoirthm

**Algorithm** *PQ-Sort*( $S$ ,  $C$ )

**Input** sequence  $S$ , comparator  $C$  for the elements of  $S$

**Output** sequence  $S$  sorted in increasing order according to  $C$

$P \leftarrow$  priority queue with comparator  $C$

**while**  $\neg S.isEmpty()$

$e \leftarrow S.removeFirst()$


$P.insert(e, \emptyset)$

**while**  $\neg P.isEmpty()$

$e \leftarrow P.removeMin().getKey()$


$S.addLast(e)$

# Heap-Sort




- Using a heap-based priority queue, the following is true:
  - Each insertion takes  $O(\log n)$  time
    - ♦ In fact the insertion of any entry  $i$  takes exactly  $O(1 + \log i)$  since this is actually the  $i^{\text{th}}$  insertion where  $1 \leq i \leq n$ .
  - Each removal takes  $O(\log n)$  time
    - ♦ In fact the removal of any entry  $i$  takes exactly  $O(1 + \log(n - i + 1))$  since this is actually the  $i^{\text{th}}$  removal (that is, some elements may have already been removed) where  $1 \leq i \leq n$ .
- Consequently, the entire algorithm consumes  $O(n \log n)$  time to have  $n$  insertions in phase I (first loop), and  $O(n \log n)$  time to have  $n$  removals in phase II (second loop) which leads to a final complexity of  $O(n \log n)$ .

# Heap-Sort



- ❑ The resulting algorithm using heap to realize the P.Q. is hence called *heap-sort*.
- ❑ Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort.



# Heap-Sort



- Consider a priority queue with  $n$  items implemented by means of a heap
  - the space used is  $O(n)$
  - methods **insert** and **removeMin** take  $O(\log n)$  time
  - methods **size**, **isEmpty**, and **min** take time  $O(1)$  time
- Using a heap-based priority queue, we can sort a sequence of  $n$  elements in  $O(n \log n)$  time
- The resulting algorithm is called heap-sort
- Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort




# Array-based Heap Implementation

- We can represent a heap with  $n$  keys by means of an array of length  $n$
- For the node at rank  $i$ 
  - the left child is at rank(index)  $2i + 1$
  - the right child is at rank  $2i + 2$
- Links between nodes are not explicitly stored
- Operation insert corresponds to inserting at rank  $n$
- Operation removeMin corresponds to removing at rank  $0$
- Yields in-place heap-sort



# Merging Two Heaps

- We are given two heaps and a key  $k$
- We create a new heap with the root node storing  $k$  and with the two heaps as subtrees
- We perform downheap to restore the heap-order property





# Bottom-up Heap Construction



- It is possible to construct a heap in  $O(n \log n)$  time.
- Simply, perform  $n$  successive insertions.
  - Since worst case for any of these insertions is  $\log n$ , the total time consumption is  $O(n \log n)$ .
- However, if all the entries (that is all the key-value pairs) are known and given in advance, we can have an alternative approach that results only on a complexity of  $O(n)$ .
- Such approach is referred to as the *bottom-up construction* approach.

# Bottom-up Heap Construction



- To simplify the description of the algorithm, let us assume that the heap is a complete binary tree; that is all levels are full.
- At each height  $i$ , we have  $2^i$  nodes
- Total number of nodes for a tree of height  $h$  is:

$$n = 2^0 + 2^1 + 2^2 + 2^3 + \dots + 2^h$$

which is equivalent to  $2^{h+1} - 1$

for example, if  $h=4$  then we have 31 nodes, which is  $2^5 - 1$

- Consequently,

- ◆  $n = 2^{h+1} - 1$



- ◆  $h = \log(n + 1) - 1$

- The construction of the heap consists of  $h + 1$  steps (which is also  $\log(n + 1)$  steps).

# Bottom-up Heap Construction



- Generally, in each step  $i$ , pairs of heaps with  $2^i - 1$  keys are merged to construct larger heaps with  $2^{i+1} - 1$  keys.
- Example: Assume  $n = 31$ ,  $h = 4$




Step # $i$	# of heaps $n+1/2^i$	# of elements in each heap $2^i - 1$ (this leads to a larger heap of $2^{i+1} - 1$ at next step $i+1$ )
1	16	1
2	8	3
3	4	7
4	2	15
5	1	31

# Bottom-up Heap Construction




- Generally, in each step  $i$ , pairs of heaps with  $2^i - 1$  keys are merged to construct larger heaps with  $2^{i+1} - 1$  keys.
- Step 1: construct  $(n + 1) / 2^1$  elementary heaps with one entry each.
- Step 2: construct  $(n + 1) / 2^2$  heaps, each storing 3 entries by joining pairs of elementary heaps and adding a new entry. New entry is added at the root, and swapping is performed if needed.



# Bottom-up Heap Construction




- :
- Step  $h + 1$  : In this last step, construct  $(n + 1) / 2^{h+1}$  heap. This is actually the final heap containing all  $n$  elements by joining the last two heaps (each storing  $(n - 1)/2$  entries) from Step  $h$ .
- As before, the new entry is added at the root, and swapping is performed if needed.
  - Note that all swap operations in each of the steps use down-heap bubbling to preserve the heap-order property.




# Bottom-up Heap Construction Example

□  $h = 3$




Step 1




Step 2

# Example (contd.)




Heap-order property is violated; we need to apply down-heap bubbling




Step 2 is concluded

# Example (contd.)




Step 3




Heap-order property was violated; Heap after applying down-heap bubbling




# Example (end)



Step 4





Final heap after applying down-heap bubbling



# Analysis


- Let us consider any internal node  $v$ . The construction of the tree rooted at  $v$  (subtree actually, except when  $v$  is the root) is proportional to the size of the paths of this tree.
- Which paths are associated with a node?
- An *associated path*,  $p(v)$ , with a node  $v$  is the one that goes to the right child of  $v$  then goes left downward until it reaches an external node (note that this path may differ from the actual downheap path).






# Analysis

- Clearly the size (number of nodes) of  $p(v)$  is equal to the: height of the tree rooted at  $v + 1$ .
- Each node in the tree belongs to at most 2 associated paths: the one rooted at  $v$  and the one coming down from its parent (note that the root and all the nodes on the left-most root-to-leaf path belong only to one associated path).





# Analysis

- Therefore, the total sizes of the paths associated with the internal nodes of the tree is at most  $2n - 1$ .
- Consequently, the construction of the tree (all its paths) using bottom-up heap construction runs in  $O(n)$  time.
- ➔ Bottom-up heap construction (while does not change the general complexity of heap-sort), it speeds up the first phase of the algorithm.

